

## ✓ Diabetes & adult income Dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

### Observations & Answers

1 Which columns in the dataset had missing values? How did you handle them?

Diabetes Dataset: Numeric columns had missing values, which were filled with mean values. Adult Income Dataset: Some rows had missing values, which were dropped.

2 Which categorical columns did you identify? How did you encode them?

Adult Income Dataset: Categorical columns like workclass, education, marital-status, etc. were encoded using LabelEncoder().

3 Difference Between Min-Max Scaling & Standardization

Min-Max Scaling:

Scales data between a fixed range (e.g., [0,1]). Preserves the shape of original distribution. Used when data does not follow a normal distribution. Standardization:

Transforms data to have zero mean and unit variance. Useful when data follows a normal (Gaussian) distribution. Preferred for algorithms like SVM, PCA, and linear regression.

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder, OneHotEncoder
import seaborn as sns
import matplotlib.pyplot as plt

diabetes = pd.read_csv("/content/drive/MyDrive/MLlab dataset/Dataset of Diabetes .csv") # Update filename accordingly

# Load Adult Income dataset
adult_income = pd.read_csv("/content/drive/MyDrive/MLlab dataset/housing.csv") # Update filename accordingly

# 1. Data Cleaning
## Identify numeric and categorical columns
numeric_cols = diabetes.select_dtypes(include=['number']).columns
categorical_cols = diabetes.select_dtypes(include=['object']).columns

## Handling Missing Values
diabetes[numeric_cols] = diabetes[numeric_cols].fillna(diabetes[numeric_cols].mean())
diabetes[categorical_cols] = diabetes[categorical_cols].fillna(diabetes[categorical_cols].mode().iloc[0])
adult_income.dropna(inplace=True) # Drop missing values in Adult dataset

## Handling Categorical Data
categorical_columns = adult_income.select_dtypes(include=['object']).columns
label_encoders = {}
for col in categorical_columns:
    le = LabelEncoder()
    adult_income[col] = le.fit_transform(adult_income[col])
    label_encoders[col] = le # Store encoder for inverse transform if needed

## Handling Outliers using IQR method
def remove_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    df[column] = np.where(df[column] < lower_bound, lower_bound, df[column])
    df[column] = np.where(df[column] > upper_bound, upper_bound, df[column])
    return df[column] # Return only the processed column

# Apply outlier removal only to numeric columns
for col in numeric_cols:
    diabetes[col] = remove_outliers(diabetes, col)

# 2. Data Transformations
## Min-Max Scaling
minmax_scaler = MinMaxScaler()
```

```


diabetes_scaled = pd.DataFrame(minmax_scaler.fit_transform(diabetes[numeric_cols]), columns=numeric_cols)
adult_income_scaled = pd.DataFrame(minmax_scaler.fit_transform(adult_income), columns=adult_income.columns)

## Standardization
standard_scaler = StandardScaler()
diabetes_standardized = pd.DataFrame(standard_scaler.fit_transform(diabetes[numeric_cols]), columns=numeric_cols)
adult_income_standardized = pd.DataFrame(standard_scaler.fit_transform(adult_income), columns=adult_income.columns)

# Save processed datasets
diabetes_scaled.to_csv("diabetes_preprocessed.csv", index=False)
adult_income_scaled.to_csv("adult_income_preprocessed.csv", index=False)

print("Data preprocessing completed.")

```

 Data preprocessing completed.

## ✓ Show Missing Values Before & After Handling

```


# Before handling missing values
print("Missing values before handling:")
print("Diabetes Dataset:\n", diabetes.isnull().sum())
print("\nAdult Income Dataset:\n", adult_income.isnull().sum())

```

```

# After handling missing values
print("\nMissing values after handling:")
print("Diabetes Dataset:\n", diabetes.isnull().sum())
print("\nAdult Income Dataset:\n", adult_income.isnull().sum())

```

 VLDL        0  
BMI        0  
CLASS     0  
dtype: int64

Adult Income Dataset:

age	0
workclass	0
fnlwgt	0
education	0
educational-num	0
marital-status	0
occupation	0
relationship	0
race	0
gender	0
capital-gain	0
capital-loss	0
hours-per-week	0
native-country	0
income	0

dtype: int64

Missing values after handling:

Diabetes Dataset:

ID	0
No_Pation	0
Gender	0
AGE	0
Urea	0
Cr	0
HbA1c	0
Chol	0
TG	0
HDL	0
LDL	0

```
relationship      0
race              0
gender            0
capital-gain      0
capital-loss      0
hours-per-week    0
native-country    0
income            0
dtype: int64
```


▼ Show Categorical Encoding

 **Generate**



Close

```
print("Categorical Columns in Adult Income Dataset:\n", categorical_columns)
print("\nEncoded Example:\n", adult_income.head())
```

 Categorical Columns in Adult Income Dataset:  
Index(['workclass', 'education', 'marital-status', 'occupation',  
 'relationship', 'race', 'gender', 'native-country', 'income'],  
 dtype='object')

Encoded Example:

	age	workclass	fnlwgt	education	educational-num	marital-status	\
0	25	4	226802	1	7	4	
1	38	4	89814	11	9	2	
2	28	2	336951	7	12	2	
3	44	4	160323	15	10	2	
4	18	0	103497	15	10	4	

	occupation	relationship	race	gender	capital-gain	capital-loss	\
0	7	3	2	1	0	0	
1	5	0	4	1	0	0	
2	11	0	4	1	0	0	
3	7	0	2	1	7688	0	
4	0	3	4	0	0	0	


  

	hours-per-week	native-country	income
0	40	39	0
1	50	39	0
2	40	39	1
3	40	39	1
4	30	39	0

▼ Before & After Scaling (Min-Max & Standardization)

```
print("Before Min-Max Scaling:\n", diabetes[numeric_cols].head())
print("\nAfter Min-Max Scaling:\n", diabetes_scaled.head())
```

```
print("\nBefore Standardization:\n", diabetes[numeric_cols].head())
print("\nAfter Standardization:\n", diabetes_standardized.head())
```

 Before Min-Max Scaling:

	ID	No_Patien	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI
0	502.0	17975.0	50.0	4.7	46.0	4.9	4.2	0.9	1.9	1.4	0.5	24.0
1	735.0	34221.0	39.0	4.5	62.0	4.9	3.7	1.4	1.1	2.1	0.6	23.0
2	420.0	47975.0	50.0	4.7	46.0	4.9	4.2	0.9	1.9	1.4	0.5	24.0
3	680.0	77365.0	50.0	4.7	46.0	4.9	4.2	0.9	1.9	1.4	0.5	24.0
4	504.0	34223.0	39.0	7.1	46.0	4.9	4.9	1.0	0.8	2.0	0.4	21.0

After Min-Max Scaling:

	ID	No_Patien	AGE	Urea	Cr	HbA1c	Chol	TG	\
0	0.627034	0.231118	0.34375	0.500	0.355	0.266892	0.406250	0.127660	
1	0.918648	0.441444	0.00000	0.475	0.515	0.266892	0.328125	0.234043	
2	0.524406	0.619508	0.34375	0.500	0.355	0.266892	0.406250	0.127660	
3	0.849812	1.000000	0.34375	0.500	0.355	0.266892	0.406250	0.127660	
4	0.629537	0.441470	0.00000	0.800	0.355	0.266892	0.515625	0.148936	

	HDL	LDL	VLDL	BMI
0	1.0000	0.209524	0.153846	0.204082
1	0.5000	0.342857	0.192308	0.163265
2	1.0000	0.209524	0.153846	0.204082
3	1.0000	0.209524	0.153846	0.204082
4	0.3125	0.323810	0.115385	0.081633

Before Standardization:

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI
0	502.0	17975.0	50.0	4.7	46.0	4.9	4.2	0.9	1.9	1.4	0.5	24.0
1	735.0	34221.0	39.0	4.5	62.0	4.9	3.7	1.4	1.1	2.1	0.6	23.0
2	420.0	47975.0	50.0	4.7	46.0	4.9	4.2	0.9	1.9	1.4	0.5	24.0
3	680.0	77365.0	50.0	4.7	46.0	4.9	4.2	0.9	1.9	1.4	0.5	24.0
4	504.0	34223.0	39.0	7.1	46.0	4.9	4.9	1.0	0.8	2.0	0.4	21.0

After Standardization:

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	\
0	0.672140	-0.919118	-0.541555	-0.074031	-0.805658	-1.335842	-0.532005	
1	1.641852	-0.087690	-2.036062	-0.190760	-0.017005	-1.335842	-0.945425	
2	0.330868	0.616204	-0.541555	-0.074031	-0.805658	-1.335842	-0.532005	
3	1.412950	2.120307	-0.541555	-0.074031	-0.805658	-1.335842	-0.532005	
4	0.680463	-0.087588	-2.036062	1.326714	-0.805658	-1.335842	0.046783	

	TG	HDL	LDL	VLDL	BMI
0	-1.200206	2.174317	-1.146921	-1.021165	-1.130562
1	-0.765541	-0.121234	-0.473190	-0.861708	-1.333654
2	-1.200206	2.174317	-1.146921	-1.021165	-1.130562
3	-1.200206	2.174317	-1.146921	-1.021165	-1.130562
4	-1.113273	-0.982065	-0.569437	-1.180623	-1.739836


## ✓ Plot Outliers Before & After Removal

```
# Boxplot before outlier handling
plt.figure(figsize=(12,6))
sns.boxplot(data=diabetes[numeric_cols])
plt.title("Before Outlier Handling")
plt.show()
```


```
# Boxplot after outlier handling
plt.figure(figsize=(12,6))
sns.boxplot(data=diabetes[numeric_cols])
plt.title("After Outlier Handling")
plt.show()
```



Before Outlier Handling



After Outlier Handling



## Save & Load Processed Data

Generate

randomly select 5 items from a list



Close

```
# Save the preprocessed data
diabetes_scaled.to_csv("diabetes_preprocessed.csv", index=False)
adult_income_scaled.to_csv("adult_income_preprocessed.csv", index=False)

# Load & display a preview
df_diabetes_processed = pd.read_csv("diabetes_preprocessed.csv")
df_adult_processed = pd.read_csv("adult_income_preprocessed.csv")

print("Preview of Preprocessed Diabetes Data:\n", df_diabetes_processed.head())
print("\nPreview of Preprocessed Adult Income Data:\n", df_adult_processed.head())
```



Preview of Preprocessed Diabetes Data:

	ID	No_Patient	AGE	Urea	Cr	HbA1c	Chol	TG	\
0	0.627034	0.231118	0.34375	0.500	0.355	0.266892	0.406250	0.127660	
1	0.918648	0.441444	0.00000	0.475	0.515	0.266892	0.328125	0.234043	
2	0.524406	0.619508	0.34375	0.500	0.355	0.266892	0.406250	0.127660	
3	0.849812	1.000000	0.34375	0.500	0.355	0.266892	0.406250	0.127660	
4	0.629537	0.441470	0.00000	0.800	0.355	0.266892	0.515625	0.148936	

	HDL	LDL	VLDL	BMI
0	1.0000	0.209524	0.153846	0.204082
1	0.5000	0.342857	0.192308	0.163265
2	1.0000	0.209524	0.153846	0.204082
3	1.0000	0.209524	0.153846	0.204082
4	0.3125	0.323810	0.115385	0.081633

Preview of Preprocessed Adult Income Data:

	age	workclass	fnlwgt	education	educational-num	marital-status	\
0	0.109589	0.50	0.145129	0.066667	0.400000	0.666667	
1	0.287671	0.50	0.052451	0.733333	0.533333	0.333333	
2	0.150685	0.25	0.219649	0.466667	0.733333	0.333333	
3	0.369863	0.50	0.100153	1.000000	0.600000	0.333333	
4	0.013699	0.00	0.061708	1.000000	0.600000	0.666667	

	occupation	relationship	race	gender	capital-gain	capital-loss	\
0	0.500000		0.6	0.5	1.0	0.000000	0.0
1	0.357143		0.0	1.0	1.0	0.000000	0.0
2	0.785714		0.0	1.0	1.0	0.000000	0.0
3	0.500000		0.0	0.5	1.0	0.076881	0.0
4	0.000000		0.6	1.0	0.0	0.000000	0.0

	hours-per-week	native-country	income
0	0.397959	0.95122	0.0
1	0.500000	0.95122	0.0
2	0.397959	0.95122	1.0
3	0.397959	0.95122	1.0
4	0.397959	0.95122	1.0

## 1. Data Cleaning

Handling Missing Values:

Diabetes Dataset:

Numeric columns: Missing values are replaced with their mean using `.fillna(diabetes[numeric_cols].mean())`.

Categorical columns: Missing values are replaced with their mode (most frequent value).

Adult Income Dataset:

All rows with missing values are dropped using `adult_income.dropna(inplace=True)`.

Handling Categorical Data:

Used `LabelEncoder()` to convert categorical columns in the Adult Income dataset into numerical values.

Handling Outliers:

Used the IQR (Interquartile Range) method to detect and replace outliers with boundary values in numeric columns of the Diabetes dataset.

## 2. Data Transformations

Min-Max Scaling (Normalization)

Applied to numeric columns in both datasets using `MinMaxScaler()`.

Standardization (Z-score scaling)

Applied to numeric columns in both datasets using `StandardScaler()`.

## ✓ Housing

```
import pandas as pd
import os

# Define file path
file_path = "/content/drive/MyDrive/MLlab dataset/housing.csv"

# Check if the file exists before loading
if os.path.exists(file_path):
    df = pd.read_csv(file_path)
    print("Loaded dataset from CSV file.")
else:
```

```

print("File not found! Check the file path.")

# Display information of all columns
print("\nDataset Info:")
print(df.info())

# Display statistical information of all numerical columns
print("\nStatistical Summary:")
print(df.describe())

# Check if "Ocean Proximity" exists (correcting column name case)
column_name = "ocean_proximity" if "ocean_proximity" in df.columns else "Ocean Proximity"

# Display the count of unique labels for "Ocean Proximity" column
if column_name in df.columns:
    print("\nOcean Proximity Value Counts:")
    print(df[column_name].value_counts())
else:
    print("\n'Ocean Proximity' column not found!")

# Display which attributes (columns) in a dataset have missing values count greater than zero
missing_values = df.isnull().sum()[df.isnull().sum() > 0]

print("\nMissing Values in Dataset:")
if missing_values.empty:
    print("No missing values found.")
else:
    print(missing_values)

```

➦ Loaded dataset from CSV file.

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 20640 entries, 0 to 20639

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	longitude	20640 non-null	float64
1	latitude	20640 non-null	float64
2	housing_median_age	20640 non-null	float64
3	total_rooms	20640 non-null	float64
4	total_bedrooms	20433 non-null	float64
5	population	20640 non-null	float64
6	households	20640 non-null	float64
7	median_income	20640 non-null	float64
8	median_house_value	20640 non-null	float64
9	ocean_proximity	20640 non-null	object

dtypes: float64(9), object(1)

memory usage: 1.6+ MB

None

Statistical Summary:

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	